

DS6372 Project 1

Carl Keusseyan, John Olan, Feby Thomas

1/31/2021

MAIN OBJECTIVES

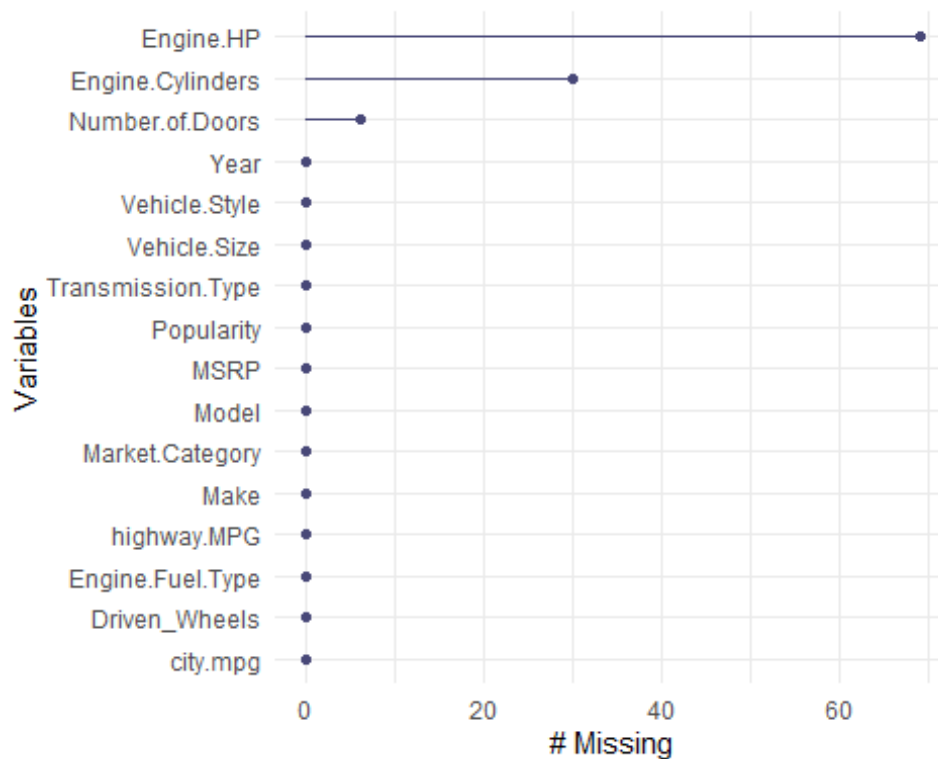
1. Read in the data
2. EDA
 - a. Deal with missing data - analyze and document how we deal with it
 - b. Examine the data - boxplots, Min, Max, look for outliers
 - c. deal with columns that are Factors - turn them into levels
 - d. Scatter plot matrix to see if there is a relationship between any 2 variables This gives us an idea on the model
3. Fit a preliminary linear model to the entire data set.
4. If residual plot shows non-constant variance, use weighted linear regression
5. Split the data into 80% train, 10% test and 10% validation First go for a simple model (no interaction terms or quadratics) using the train set, predict and compare to the test set, and then validate.
6. Fit the model using Single tree, prune as it may.
7. Fit to random forest. Grid search parameters.
8. Fit using lasso regularization. Grid search parameters.
9. Tabulate all the Test RMSEs to see which is lowest.

The codes as adequately clearly commented upon.

EXPLORATORY DATA ANALYSES

- Deal with missing data - analyze and document how we deal with it

#Let us plot what variable are missing
`gg_miss_var(car.new)`



As we can see from the above figure. The Engine.HP (70) and Engine.Cylinders (30) have some missing values.

Using the View() we realized that The Market.Category contained a substantial amount of 'N/As' (3600 missing values) which were not captured by `gg_miss_var(car.new)` function.

This values in the Market Category were subjectively inputed and could not help with prediction.

- Examine the data - boxplots, Min, Max, look for outliers
- deal with columns that are Factors - turn them into levels
- Scatter plot matrix to see if there is a relationship between any 2 variables
This gives us an idea on the model

KNITTED FILE FROM R-MARKDOWN

#set Working directory

```
setwd("C:/Users/olani/OneDrive/Documents/Data Science/SMU-Data Science/Applied Statistics/Project1Details_2021/Project1Details_2021")
```

#Read in the CSV file

```
car.new<-read.csv("data1.csv")  
attach(car.new)
```

#examine the newly read in dataset car.new

```
head(car.new)
```

```
##      Make      Model Year      Engine.Fuel.Type Engine.HP Engine.Cylinders  
## 1  BMW 1 Series M 2011 premium unleaded (required)      335  
6  
## 2  BMW 1 Series 2011 premium unleaded (required)      300  
6  
## 3  BMW 1 Series 2011 premium unleaded (required)      300  
6  
## 4  BMW 1 Series 2011 premium unleaded (required)      230  
6  
## 5  BMW 1 Series 2011 premium unleaded (required)      230  
6  
## 6  BMW 1 Series 2012 premium unleaded (required)      230  
6  
##      Transmission.Type      Driven_Wheels      Number.of.Doors  
## 1      MANUAL rear wheel drive      2  
## 2      MANUAL rear wheel drive      2  
## 3      MANUAL rear wheel drive      2  
## 4      MANUAL rear wheel drive      2  
## 5      MANUAL rear wheel drive      2  
## 6      MANUAL rear wheel drive      2  
##      Market.Category      Vehicle.Size      Vehicle.Style      highway  
.MPG  
## 1 Factory Tuner,Luxury,High-Performance      Compact      Coupe  
26  
## 2      Luxury,Performance      Compact      Convertible  
28  
## 3      Luxury,High-Performance      Compact      Coupe  
28  
## 4      Luxury,Performance      Compact      Coupe  
28  
## 5      Luxury      Compact      Convertible  
28  
## 6      Luxury,Performance      Compact      Coupe  
28  
##      city.mpg      Popularity      MSRP
```

```
## 1      19      3916 46135
## 2      19      3916 40650
## 3      20      3916 36350
## 4      18      3916 29450
## 5      18      3916 34500
## 6      18      3916 31200
```

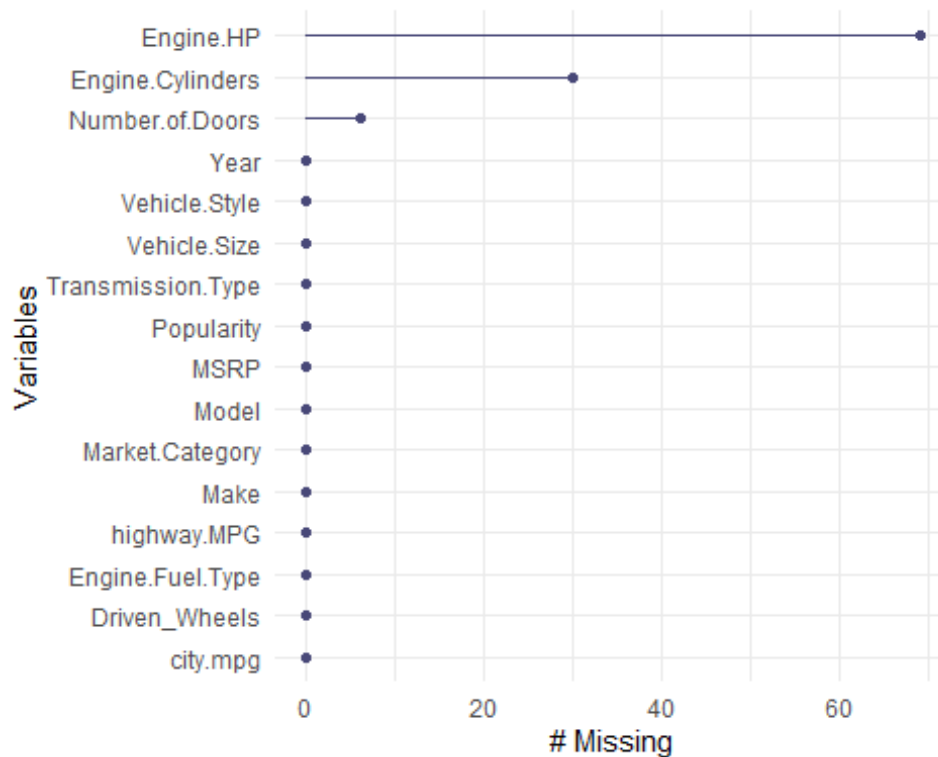
```
str(car.new)
```

```
## 'data.frame':    11914 obs. of  16 variables:
##  $ Make           : chr  "BMW" "BMW" "BMW" "BMW" ...
##  $ Model          : chr  "1 Series M" "1 Series" "1 Series" "1 Series" .
..
##  $ Year           : int   2011 2011 2011 2011 2011 2012 2012 2012 2012 20
13 ...
##  $ Engine.Fuel.Type : chr   "premium unleaded (required)" "premium unleaded
(required)" "premium unleaded (required)" "premium unleaded (required)" ...
##  $ Engine.HP       : int   335 300 300 230 230 230 300 300 230 230 ...
##  $ Engine.Cylinders : int    6 6 6 6 6 6 6 6 6 6 ...
##  $ Transmission.Type: chr   "MANUAL" "MANUAL" "MANUAL" "MANUAL" ...
##  $ Driven_Wheels   : chr   "rear wheel drive" "rear wheel drive" "rear whe
el drive" "rear wheel drive" ...
##  $ Number.of.Doors  : int    2 2 2 2 2 2 2 2 2 2 ...
##  $ Market.Category  : chr   "Factory Tuner,Luxury,High-Performance" "Luxury
,Performance" "Luxury,High-Performance" "Luxury,Performance" ...
##  $ Vehicle.Size     : chr   "Compact" "Compact" "Compact" "Compact" ...
##  $ Vehicle.Style    : chr   "Coupe" "Convertible" "Coupe" "Coupe" ...
##  $ highway.MPG      : int   26 28 28 28 28 28 26 28 28 27 ...
##  $ city.mpg         : int   19 19 20 18 18 18 17 20 18 18 ...
##  $ Popularity       : int   3916 3916 3916 3916 3916 3916 3916 3916 3916 39
16 ...
##  $ MSRP            : int   46135 40650 36350 29450 34500 31200 44100 39300
36900 37200 ...
```

```
### Deal with Missing Data
```

```
#let us plot what variable are missing
```

```
gg_miss_var(car.new)
```



```
# we can see the missing data
#   Engine HP around about 70 values
#   Engine Cylinders about 30 values
#   Number of doors about 5

# Display the number of rows and columns
dim(car.new) # we know we have 11914 rows of data with 16 columns

## [1] 11914    16

# removing the Market.category column because it has too many N/As - 3600 rows
# are missing this value, and this amount of lack of data will compromise our
# analysis and the model. At this point we are going to take it out.

car.new <- subset(car.new, select =-c(Market.Category))

str(car.new)

## 'data.frame':    11914 obs. of  15 variables:
## $ Make           : chr  "BMW" "BMW" "BMW" "BMW" ...
## $ Model          : chr  "1 Series M" "1 Series" "1 Series" "1 Series" .
## ..
## $ Year           : int   2011 2011 2011 2011 2011 2012 2012 2012 2012 20
## 13 ...
## $ Engine.Fuel.Type : chr   "premium unleaded (required)" "premium unleaded
## (required)" "premium unleaded (required)" "premium unleaded (required)" ...
## $ Engine.HP       : int   335 300 300 230 230 230 300 300 230 230 ...
```

```
## $ Engine.Cylinders : int 6 6 6 6 6 6 6 6 6 6 ...
## $ Transmission.Type: chr "MANUAL" "MANUAL" "MANUAL" "MANUAL" ...
## $ Driven_Wheels : chr "rear wheel drive" "rear wheel drive" "rear wheel drive" "rear wheel drive" ...
## $ Number.of.Doors : int 2 2 2 2 2 2 2 2 2 2 ...
## $ Vehicle.Size : chr "Compact" "Compact" "Compact" "Compact" ...
## $ Vehicle.Style : chr "Coupe" "Convertible" "Coupe" "Coupe" ...
## $ highway.MPG : int 26 28 28 28 28 28 26 28 28 27 ...
## $ city.mpg : int 19 19 20 18 18 18 17 20 18 18 ...
## $ Popularity : int 3916 3916 3916 3916 3916 3916 3916 3916 3916 3916 ...
## $ MSRP : int 46135 40650 36350 29450 34500 31200 44100 39300 36900 37200 ...
```

###Feature Engineering####convert features to factors

##Convert relevant features to factors

```
car.new1 <- car.new %>%
```

```
  mutate(Make = as.factor(Make),
         Popularity = as.factor(Popularity),
         Model = as.factor(Model),
         Vehicle.Size = as.factor(Vehicle.Size),
         Vehicle.Style = as.factor(Vehicle.Style),
         Number.of.Doors = as.factor(Number.of.Doors),
         Driven_Wheels = as.factor(Driven_Wheels),
         Transmission.Type=as.factor(Transmission.Type),
         Engine.Cylinders = as.factor(Engine.Cylinders),
         Engine.Fuel.Type = as.factor(Engine.Fuel.Type))
```

```
str(car.new1)
```

```
## 'data.frame': 11914 obs. of 15 variables:
## $ Make : Factor w/ 48 levels "Acura","Alfa Romeo",...: 6 6 6 6 6 6 6 6 ...
## $ Model : Factor w/ 915 levels "1 Series","1 Series M",...: 2 1 1 1 1 1 1 1 ...
## $ Year : int 2011 2011 2011 2011 2011 2012 2012 2012 2012 2013 ...
## $ Engine.Fuel.Type : Factor w/ 11 levels "", "diesel", "electric",...: 10 10 10 10 10 10 10 10 ...
## $ Engine.HP : int 335 300 300 230 230 230 300 300 230 230 ...
## $ Engine.Cylinders : Factor w/ 9 levels "0","3","4","5",...: 5 5 5 5 5 5 5 5 ...
## $ Transmission.Type: Factor w/ 5 levels "AUTOMATED_MANUAL",...: 4 4 4 4 4 4 4 4 ...
## $ Driven_Wheels : Factor w/ 4 levels "all wheel drive",...: 4 4 4 4 4 4 4 4 ...
## $ Number.of.Doors : Factor w/ 3 levels "2","3","4": 1 1 1 1 1 1 1 1 1 1 ...
## $ Vehicle.Size : Factor w/ 3 levels "Compact","Large",...: 1 1 1 1 1 1 1 1
```



```

1 1 1 1 ...
## $ Vehicle.Style      : Factor w/ 16 levels "2dr Hatchback",...: 9 7 9 9 7 9
7 9 7 7 ...
## $ highway.MPG        : int   26 28 28 28 28 28 26 28 28 27 ...
## $ city.mpg           : int   19 19 20 18 18 18 17 20 18 18 ...
## $ Popularity         : Factor w/ 48 levels "2","21","26",...: 47 47 47 47 47
47 47 47 47 47 ...
## $ MSRP               : int   46135 40650 36350 29450 34500 31200 44100 39300
36900 37200 ...

```

#list of level in Make. I can't merge the make to reduce the levels. They are different.

```

make <- unique(car.new1$Make)
make

```

```

## [1] BMW          Audi          FIAT          Mercedes-Benz Chrysler
## [6] Nissan        Volvo         Mazda         Mitsubishi    Ferrari
## [11] Alfa Romeo    Toyota        McLaren       Maybach       Pontiac
## [16] Porsche      Saab          GMC           Hyundai       Plymouth
## [21] Honda        Oldsmobile    Suzuki        Ford          Cadillac
## [26] Kia          Bentley       Chevrolet     Dodge         Lamborghini
## [31] Lincoln      Subaru        Volkswagen    Spyker        Buick
## [36] Acura        Rolls-Royce   Maserati      Lexus         Aston Martin
## [41] Land Rover    Lotus         Infiniti      Scion         Genesis
## [46] HUMMER       Tesla         Bugatti
## 48 Levels: Acura Alfa Romeo Aston Martin Audi Bentley BMW Bugatti ... Volvo

```

###EDA

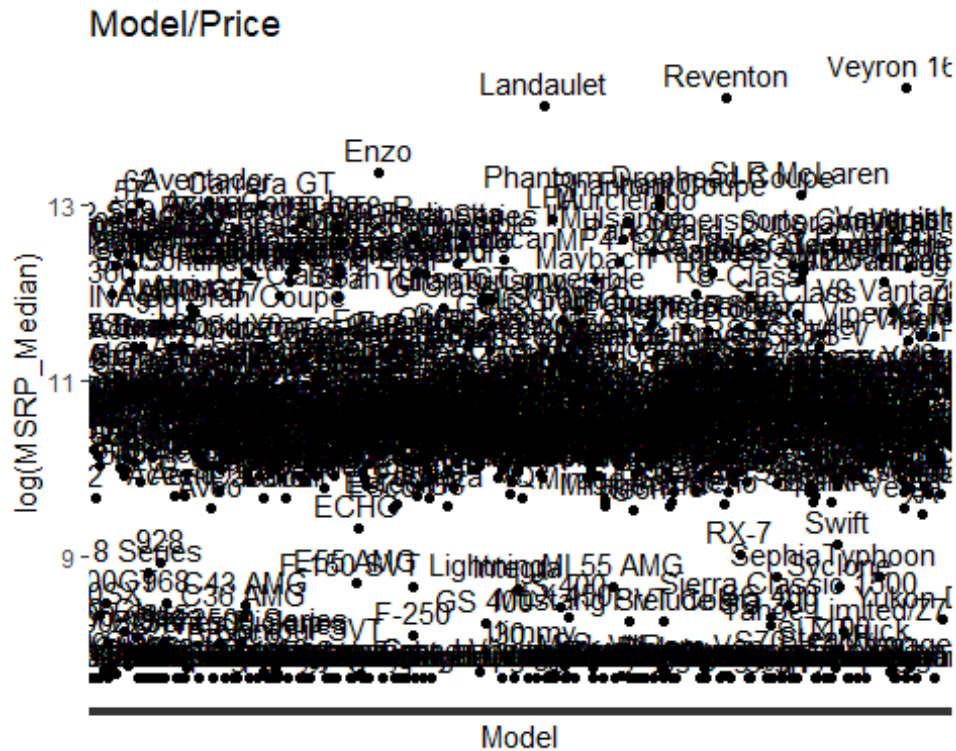
#plot some viz
#Model versus MSRP. No trend transformed or not
 ##

```

car.new1 %>% group_by(Model) %>% summarise(MSRP_Median = median(MSRP)) %>%
  arrange(desc(MSRP_Median)) %>%
  ggplot(aes(x=Model, y = log(MSRP_Median))) +
  geom_point() +
  geom_text(aes(label=Model),hjust=0.5, vjust=-0.5) + #this line adds label to
o the datapoints so I can where the outliers come from
  theme(axis.text.x = element_blank()) +
  xlab("Model") +
  ggtitle("Model/Price")

## `summarise()` ungrouping output (override with `.groups` argument)

```




```
#Make and MSRP
#High Leverage from the very expensive vehicles
#remove or keep @ MSRP> USD$100,000?

#The primary objective for this EDA section is to identify the variables that
have association with MSRP (the response)
#The code below serves as template. So I just changed out the variables as ne
eded, to view their relationship.

car.new1 %>% group_by(Make) %>% summarise(MSRP_Median = median(MSRP)) %>%
  arrange(desc(MSRP_Median)) %>%
  ggplot(aes(x = Make, y = (MSRP_Median))) +
  geom_point() +
  geom_text(aes(label=Make),hjust=0.5, vjust=-0.5) + #this line adds label to
the datapoints so I can where the outliers come from
  theme(axis.text.x = element_blank()) +
  xlab("Make") +
  ggtitle("Make/Price")

## `summarise()` ungrouping output (override with `.groups` argument)
```



###Post-EDA 1

#I subset out the highly expensive cars that have high Leverage and are 'outliers'.

```
car_nonEx <- car.new %>%
  filter(MSRP < 100000)
```

###exclude missing values

```
car_nonEx <- car_nonEx[complete.cases(car_nonEx),] #exclude NA and missing data
```

###Feature engineering and EDA 2

```
car_nonEx <- car_nonEx %>%
  mutate(Make = as.factor(Make),
         Popularity = as.factor(Popularity),
         Model = as.factor(Model),
         Vehicle.Size = as.factor(Vehicle.Size),
         Vehicle.Style = as.factor(Vehicle.Style),
         Number.of.Doors = as.factor(Number.of.Doors),
         Driven_Wheels = as.factor(Driven_Wheels),
         Transmission.Type = as.factor(Transmission.Type),
         Engine.Cylinders = as.factor(Engine.Cylinders),
         Engine.Fuel.Type = as.factor(Engine.Fuel.Type))
```

*#Vehicle style has some levels that can be merged without losing info.
#Below is to replace "4Dr SUV" and "2Dr SUV" with "SUV", given that 4Dr and 2
Dr are already captured in Number.of.Doors column.*

```
car_nonEx.new <- car_nonEx %>%  
  mutate(Veh.Style_recode1 = Vehicle.Style)
```

```
car_nonEx.new$Veh.Style_recode = str_replace_all(as.character(car_nonEx.new$Veh.Style_recode), "2dr SUV", "SUV")  
car_nonEx.new$Veh.Style_recode = str_replace_all(as.character(car_nonEx.new$Veh.Style_recode), "4dr SUV", "SUV")
```

```
car_nonEx.new$Veh.Style_recode = as.factor(car_nonEx.new$Veh.Style_recode)
```

#how does it look now.

```
str(car_nonEx.new)
```

```
## 'data.frame': 11182 obs. of 17 variables:  
## $ Make : Factor w/ 39 levels "Acura","Alfa Romeo",...: 5 5 5 5  
5 5 5 5 5 5 ...  
## $ Model : Factor w/ 811 levels "1 Series","1 Series M",...: 2 1  
1 1 1 1 1 1 1 1 ...  
## $ Year : int 2011 2011 2011 2011 2011 2012 2012 2012 2012 20  
13 ...  
## $ Engine.Fuel.Type : Factor w/ 10 levels "", "diesel", "electric",...: 9 9 9  
9 9 9 9 9 9 9 ...  
## $ Engine.HP : int 335 300 300 230 230 230 300 300 230 230 ...  
## $ Engine.Cylinders : Factor w/ 8 levels "0","3","4","5",...: 5 5 5 5 5 5 5 5  
5 5 5 ...  
## $ Transmission.Type: Factor w/ 5 levels "AUTOMATED_MANUAL",...: 4 4 4 4 4  
4 4 4 4 4 ...  
## $ Driven_Wheels : Factor w/ 4 levels "all wheel drive",...: 4 4 4 4 4 4  
4 4 4 4 ...  
## $ Number.of.Doors : Factor w/ 3 levels "2","3","4": 1 1 1 1 1 1 1 1 1 1  
...  
## $ Vehicle.Size : Factor w/ 3 levels "Compact","Large",...: 1 1 1 1 1 1  
1 1 1 1 ...  
## $ Vehicle.Style : Factor w/ 16 levels "2dr Hatchback",...: 9 7 9 9 7 9  
7 9 7 7 ...  
## $ highway.MPG : int 26 28 28 28 28 28 26 28 28 27 ...  
## $ city.mpg : int 19 19 20 18 18 18 17 20 18 18 ...  
## $ Popularity : Factor w/ 39 levels "21","26","61",...: 38 38 38 38 3  
8 38 38 38 38 38 ...  
## $ MSRP : int 46135 40650 36350 29450 34500 31200 44100 39300  
36900 37200 ...  
## $ Veh.Style_recode1: Factor w/ 16 levels "2dr Hatchback",...: 9 7 9 9 7 9
```

```


7 9 7 7 ...
## $ Veh.Style_recode : Factor w/ 15 levels "2dr Hatchback",...: 7 5 7 7 5 7
5 7 5 5 ...

#View result

#reorder argument with arrange the levels in ascending order, which makes it
easy to view the trend.
#This is also a utility code chunk that I was adapting for different plots of
different variables.

car_nonEx %>%
  #group_by(Make) %>% summarise(MSRP_Median = median(MSRP)) %>%
  #arrange(desc(MSRP_Median)) %>%
  #ggplot(aes(x = reorder(Make, -MSRP_Median), y = MSRP_Median, fill = Engine.
Cylinders)) +
  ggplot(aes(x = reorder(Make, -MSRP), y = MSRP, fill = as.factor(Engine.Cylin
ders))) +
  geom_bar(stat="identity") +
  #geom_text(aes(label=Make), hjust=0.5, vjust=-0.5) +
  #theme(axis.text.x = element_blank(angle = 45)) +
  xlab("Make") +
  ggtitle("Make/Price")

```



```

#Vehicle.Style had 16 variables with some that are redundant.
car_nonEx.new %>%

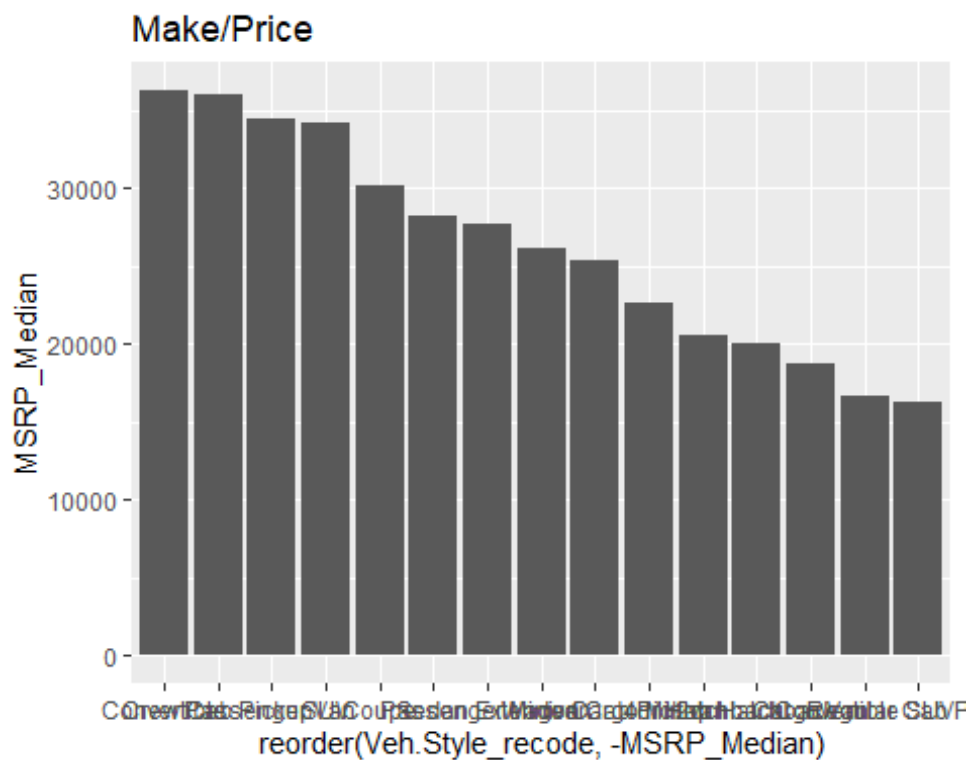
```

```

group_by(Veh.Style_recode) %>% summarise(MSRP_Median = median(MSRP)) %>%
#arrange(desc(MSRP_Median)) %>%
#ggplot(aes(x = reorder(Make,-MSRP_Median), y = MSRP_Median, fill = Engine.
Cylinders)) +
ggplot(aes(x = reorder(Veh.Style_recode,-MSRP_Median), y = MSRP_Median)) +
geom_bar(stat="identity") +
#geom_text(aes(label=Make),hjust=0.5, vjust=-0.5) + #this line adds label t
o the datapoints so I can where the outliers come from
#theme(axis.text.x = element_blank(angle = 45)) +
#xlab("Make") +
ggtitle("Make/Price")

## `summarise()` ungrouping output (override with `.groups` argument)

```



```

#View names of Levels that makes up the factor
unique(car_nonEx.new$Veh.Style_recode)

## [1] Coupe          Convertible      Sedan
## [4] Wagon          4dr Hatchback   2dr Hatchback
## [7] SUV            Passenger Minivan Cargo Minivan
## [10] Crew Cab Pickup Regular Cab Pickup Extended Cab Pickup
## [13] Cargo Van      Convertible SUV Passenger Van
## 15 Levels: 2dr Hatchback 4dr Hatchback Cargo Minivan Cargo Van ... Wagon

colSums(is.na(car_nonEx.new))

##           Make           Model           Year Engine.Fuel.Type
##           0             0             0           0

```

```
##      Engine.HP  Engine.Cylinders  Transmission.Type    Driven_Wheels
##           0           0           0           0
##  Number.of.Doors    Vehicle.Size    Vehicle.Style    highway.MPG
##           0           0           0           0
##      city.mpg    Popularity    MSRP  Veh.Style_recode1
##           0           0           0           0
##  Veh.Style_recode
##           0
```

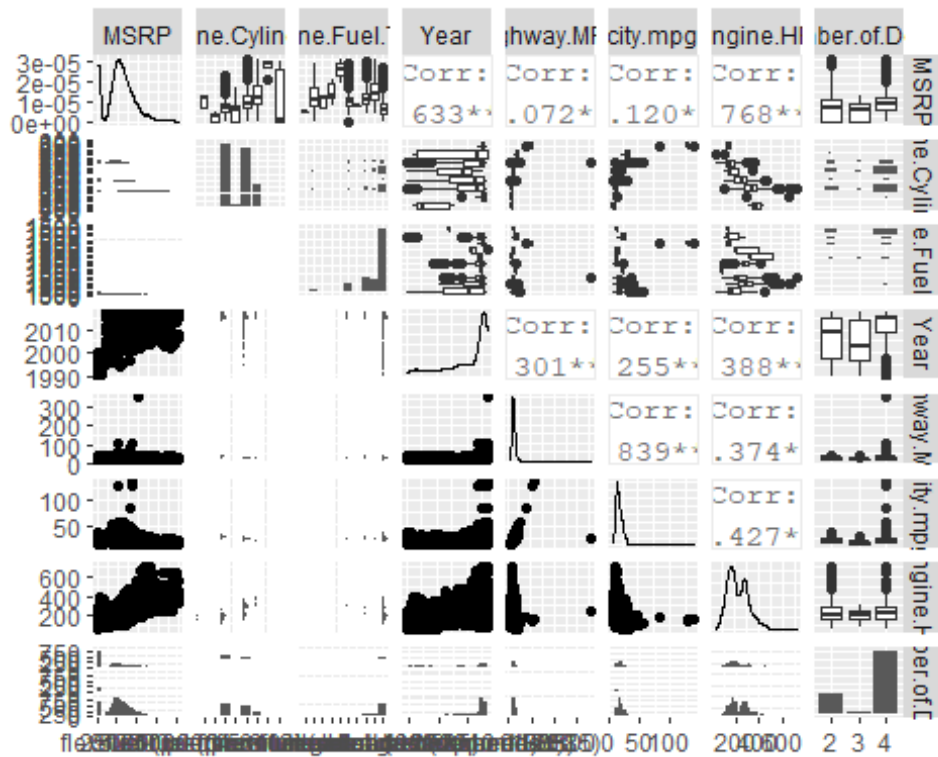
##EDA-3

#For a pairwise view of plots
#My plan was to view a set 5 or 6 features at a time instead of the whole thing.
#I played with log, squared of the variables as I.
#these plots show that city.MPG and Highway MPG do not show any considerable association with MSRP.
#Make and model do not have meaningful association with the MSRP below 100000
.

```
library(GGally)
#Pairwise plots
df2.new <- car_nonEx.new %>%
  mutate(log.MSRP = log(MSRP),
         sq.Engine.HP=(Engine.HP)^2) %>%
  dplyr::select(c(MSRP, Driven_Wheels, Transmission.Type, Number.of.Doors,
Vehicle.Size, Veh.Style_recode))

df2.raw <- car_nonEx %>%
  dplyr::select(c(MSRP, Engine.Cylinders, Engine.Fuel.Type, Year,highwa
y.MPG, city.mpg, Engine.HP, Number.of.Doors))

pairs(df2.raw)
```

#what levels are in fuel type

```
unique(car_nonEx$Engine.Cylinders)
```

```
## [1] 6 4 5 8 12 0 3 10
```

```
## Levels: 0 3 4 5 6 8 10 12
```

```
str(car_nonEx)
```

```
## 'data.frame': 11182 obs. of 15 variables:
```

```
## $ Make : Factor w/ 39 levels "Acura","Alfa Romeo",...: 5 5 5 5 5 5 5 5 ...
```

```
## $ Model : Factor w/ 811 levels "1 Series","1 Series M",...: 2 1 1 1 1 1 1 1 ...
```

```
## $ Year : int 2011 2011 2011 2011 2011 2012 2012 2012 2012 2013 ...
```

```
## $ Engine.Fuel.Type : Factor w/ 10 levels "", "diesel", "electric",...: 9 9 9 9 9 9 9 9 ...
```

```
## $ Engine.HP : int 335 300 300 230 230 230 300 300 230 230 ...
```

```
## $ Engine.Cylinders : Factor w/ 8 levels "0", "3", "4", "5",...: 5 5 5 5 5 5 5 5 ...
```

```
## $ Transmission.Type: Factor w/ 5 levels "AUTOMATED_MANUAL",...: 4 4 4 4 4 4 4 4 ...
```

```
## $ Driven_Wheels : Factor w/ 4 levels "all wheel drive",...: 4 4 4 4 4 4 4 4 ...
```

```
## $ Number.of.Doors : Factor w/ 3 levels "2", "3", "4": 1 1 1 1 1 1 1 1 1 1 ...
```

```
## $ Vehicle.Size : Factor w/ 3 levels "Compact", "Large",...: 1 1 1 1 1 1 1 1
```

```
1 1 1 1 ...
## $ Vehicle.Style      : Factor w/ 16 levels "2dr Hatchback",...: 9 7 9 9 7 9
7 9 7 7 ...
## $ highway.MPG        : int   26 28 28 28 28 28 26 28 28 27 ...
## $ city.mpg           : int   19 19 20 18 18 18 17 20 18 18 ...
## $ Popularity         : Factor w/ 39 levels "21","26","61",...: 38 38 38 38 3
8 38 38 38 38 38 ...
## $ MSRP               : int   46135 40650 36350 29450 34500 31200 44100 39300
36900 37200 ...
```

```
view(car_nonEx$Engine.Fuel.Type)
```

```
#Grouped plot
```

```
car_nonEx %>%
```

```
group_by(Engine.Fuel.Type) %>%
```

```
#summarise(MSRP_Median = median(MSRP)) %>%
```

```
#arrange(desc(MSRP_Median)) %>%
```

```
ggplot(aes(x = reorder(Engine.Fuel.Type, -MSRP), y = MSRP)) +
```

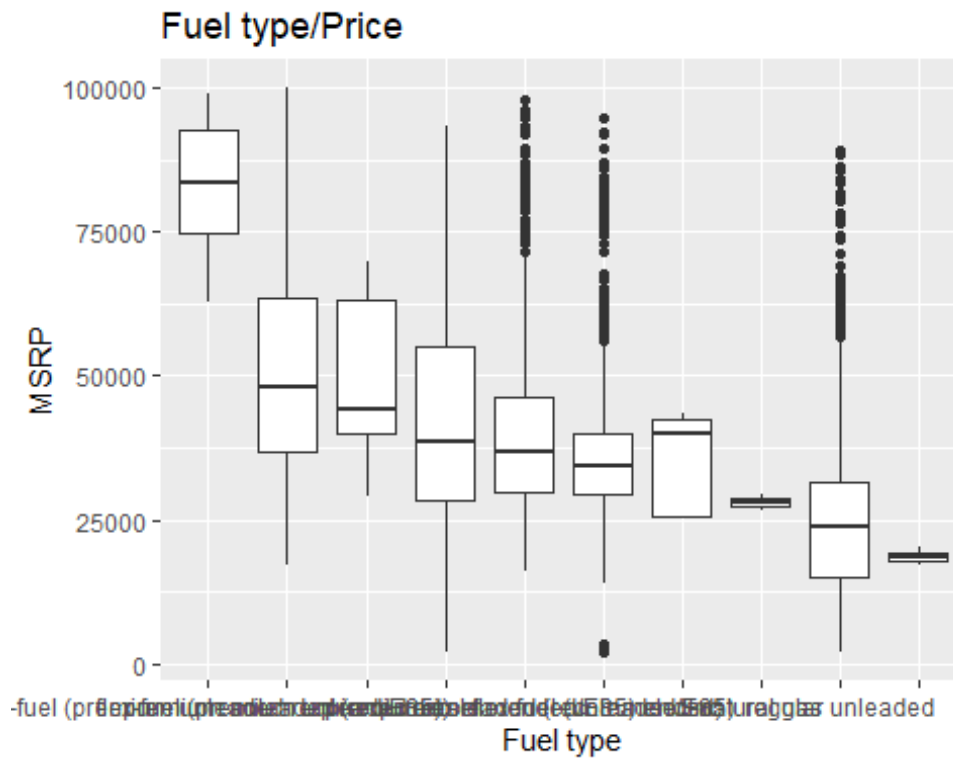
geom_boxplot() +

```
#geom_text(aes(label=Make),hjust=0.5, vjust=-0.5) + #this line adds label t
o the datapoints so I can where the outliers come from
```

```
#theme(axis.text.x = element_blank(angle = 45)) +
```

```
xlab("Fuel type") +
```


```
ggtitle("Fuel type/Price")
```



#correlation of the selected variables

```
df2.raw %>% ggcorr(palette = "RdBu", label = TRUE, hjust= 0.9, layout.exp = 1.2, name = "Spearman correlation coeff. ( $\rho$ )")

## Warning in ggcorr(., palette = "RdBu", label = TRUE, hjust = 0.9, layout.e
xp
## = 1.2, : data in column(s) 'Engine.Cylinders', 'Engine.Fuel.Type',
## 'Number.of.Doors' are not numeric and were ignored
```




#No association with City MPG/highway MPG so drop from model.

#MSRP histogram to see MSRP distribution. There was some very very cheap cars 0 to 5000 that are seriously skewing the result.

```
car_nonEx.new %>%
  ggplot(aes(x =MSRP)) +
  geom_histogram() +
  #xlab("Fuel type") +
  ggtitle("MSRP")
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
car_v.cheap = car_nonEx.new %>% filter(MSRP<5000)  
car_nonEx.new = car_nonEx.new %>% filter(MSRP >=5000)
```

There is no association of MSRP with City MPG nor Highway MPG

```
plot( MSRP~(highway.MPG), data = car_nonEx)
```



##Multiple linear regression

##Fit a preliminary model using the variables that have some association with MSRP

```
Model1<- lm(log(MSRP)~Year+(Engine.HP)^2+Transmission.Type+Driven_Wheels+Numb
er.of.Doors+Vehicle.Size+Veh.Style_recode+Engine.Fuel.Type, data=car_nonEx.ne
w)
summary(Model1)
```

```
##
## Call:
## lm(formula = log(MSRP) ~ Year + (Engine.HP)^2 + Transmission.Type +
##   Driven_Wheels + Number.of.Doors + Vehicle.Size + Veh.Style_recode +
##   Engine.Fuel.Type, data = car_nonEx.new)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8284 -0.1124 -0.0061  0.1092  1.2706
##
## Coefficients:
##                                     Estimate
## (Intercept)                      -2.109e+01
## Year                                1.526e-02
## Engine.HP                          2.481e-03
## Transmission.TypeAUTOMATIC         -2.049e-02
```

## Transmission.TypeDIRECT_DRIVE	5.550e-02	
## Transmission.TypeMANUAL	-1.549e-01	
## Transmission.TypeUNKNOWN	-1.339e+00	
## Driven_Wheelsfour wheel drive	4.778e-02	
## Driven_Wheelsfront wheel drive	-1.047e-01	
## Driven_Wheelsrear wheel drive	-4.550e-02	
## Number.of.Doors3	-2.647e-02	
## Number.of.Doors4	-5.210e-02	
## Vehicle.SizeLarge	1.405e-01	
## Vehicle.SizeMidsize	3.847e-02	
## Veh.Style_recode4dr Hatchback	2.990e-02	
## Veh.Style_recodeCargo Minivan	5.535e-02	
## Veh.Style_recodeCargo Van	-1.554e-02	
## Veh.Style_recodeConvertible	2.227e-01	
## Veh.Style_recodeConvertible SUV	1.218e-01	
## Veh.Style_recodeCoupe	4.841e-02	
## Veh.Style_recodeCrew Cab Pickup	2.844e-02	
## Veh.Style_recodeExtended Cab Pickup	-4.492e-02	
## Veh.Style_recodePassenger Minivan	1.853e-01	
## Veh.Style_recodePassenger Van	9.029e-02	
## Veh.Style_recodeRegular Cab Pickup	-1.877e-01	
## Veh.Style_recodeSedan	9.574e-02	
## Veh.Style_recodeSUV	1.562e-01	
## Veh.Style_recodeWagon	1.251e-01	
## Engine.Fuel.Typediesel	4.612e-01	
## Engine.Fuel.Typeelectric	4.157e-01	
## Engine.Fuel.Typeflex-fuel (premium unleaded recommended/E85)	3.884e-01	
## Engine.Fuel.Typeflex-fuel (premium unleaded required/E85)	4.718e-01	
## Engine.Fuel.Typeflex-fuel (unleaded/E85)	6.578e-02	
## Engine.Fuel.Typenatural gas	3.929e-01	
## Engine.Fuel.Typepremium unleaded (recommended)	2.252e-01	
## Engine.Fuel.Typepremium unleaded (required)	3.407e-01	
## Engine.Fuel.Typeregular unleaded	6.938e-02	
##	Std. Error t	
value		
## (Intercept)	9.827e-01	-2
1.465		
## Year	4.877e-04	3
1.291		
## Engine.HP	3.718e-05	6
6.718		
## Transmission.TypeAUTOMATIC	1.044e-02	-
1.962		
## Transmission.TypeDIRECT_DRIVE	1.376e-01	
0.403		
## Transmission.TypeMANUAL	1.085e-02	-1
4.278		
## Transmission.TypeUNKNOWN	1.131e-01	-1
1.834		
## Driven_Wheelsfour wheel drive	8.865e-03	

5.389		
## Driven_Wheelsfront wheel drive	6.293e-03	-1
6.632		
## Driven_Wheelsrear wheel drive	7.042e-03	-
6.461		
## Number.of.Doors3	2.570e-02	-
1.030		
## Number.of.Doors4	2.206e-02	-
2.362		
## Vehicle.SizeLarge	7.602e-03	1
8.480		
## Vehicle.SizeMidsize	5.727e-03	
6.718		
## Veh.Style_recode4dr Hatchback	2.448e-02	
1.221		
## Veh.Style_recodeCargo Minivan	3.278e-02	
1.689		
## Veh.Style_recodeCargo Van	3.875e-02	-
0.401		
## Veh.Style_recodeConvertible	1.388e-02	1
6.052		
## Veh.Style_recodeConvertible SUV	4.739e-02	
2.570		
## Veh.Style_recodeCoupe	1.354e-02	
3.576		
## Veh.Style_recodeCrew Cab Pickup	2.552e-02	
1.114		
## Veh.Style_recodeExtended Cab Pickup	2.492e-02	-
1.802		
## Veh.Style_recodePassenger Minivan	2.534e-02	
7.312		
## Veh.Style_recodePassenger Van	3.343e-02	
2.701		
## Veh.Style_recodeRegular Cab Pickup	1.858e-02	-1
0.103		
## Veh.Style_recodeSedan	2.371e-02	
4.039		
## Veh.Style_recodeSUV	2.370e-02	
6.590		
## Veh.Style_recodeWagon	2.503e-02	
4.996		
## Engine.Fuel.Typediesel	1.134e-01	
4.068		
## Engine.Fuel.Typeelectric	1.854e-01	
2.242		
## Engine.Fuel.Typeflex-fuel (premium unleaded recommended/E85)	1.185e-01	
3.278		
## Engine.Fuel.Typeflex-fuel (premium unleaded required/E85)	1.225e-01	
3.852		
## Engine.Fuel.Typeflex-fuel (unleaded/E85)	1.124e-01	

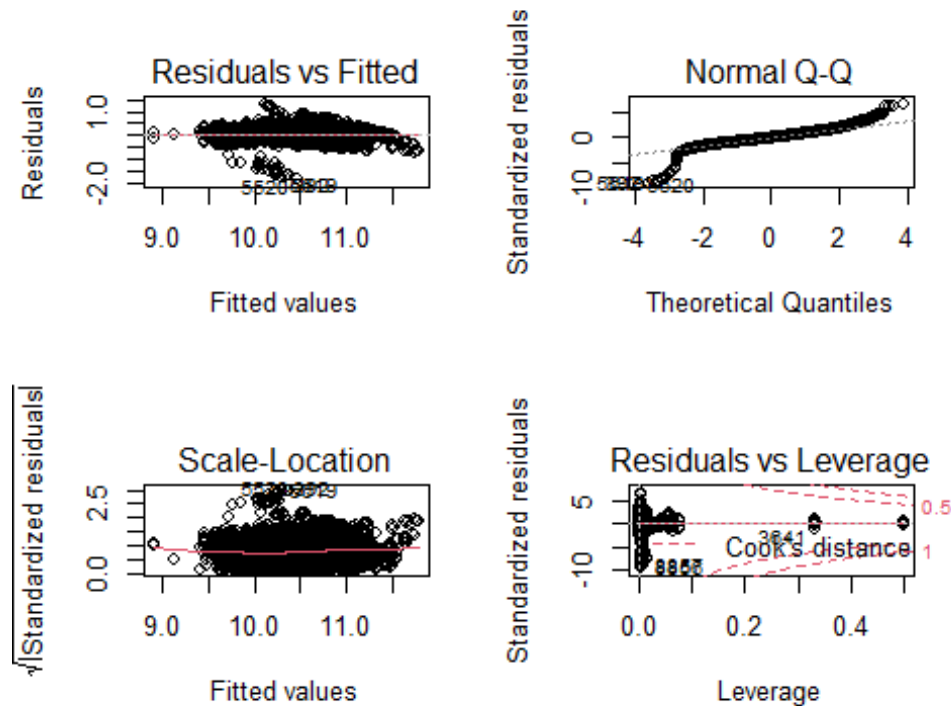
```

0.585
## Engine.Fuel.Typenatural gas 1.771e-01
2.218
## Engine.Fuel.Typepremium unleaded (recommended) 1.123e-01
2.005
## Engine.Fuel.Typepremium unleaded (required) 1.123e-01
3.034
## Engine.Fuel.Typeregular unleaded 1.121e-01
0.619
## Pr(>|t|)
## (Intercept) < 2e-16 ***
## Year < 2e-16 ***
## Engine.HP < 2e-16 ***
## Transmission.TypeAUTOMATIC 0.049734 *
## Transmission.TypeDIRECT_DRIVE 0.686600
## Transmission.TypeMANUAL < 2e-16 ***
## Transmission.TypeUNKNOWN < 2e-16 ***
## Driven_Wheelsfour wheel drive 7.24e-08 ***
## Driven_Wheelsfront wheel drive < 2e-16 ***
## Driven_Wheelsrear wheel drive 1.09e-10 ***
## Number.of.Doors3 0.303204
## Number.of.Doors4 0.018219 *
## Vehicle.SizeLarge < 2e-16 ***
## Vehicle.SizeMidsize 1.94e-11 ***
## Veh.Style_recode4dr Hatchback 0.221987
## Veh.Style_recodeCargo Minivan 0.091333 .
## Veh.Style_recodeCargo Van 0.688412
## Veh.Style_recodeConvertible < 2e-16 ***
## Veh.Style_recodeConvertible SUV 0.010198 *
## Veh.Style_recodeCoupe 0.000351 ***
## Veh.Style_recodeCrew Cab Pickup 0.265215
## Veh.Style_recodeExtended Cab Pickup 0.071513 .
## Veh.Style_recodePassenger Minivan 2.85e-13 ***
## Veh.Style_recodePassenger Van 0.006923 **
## Veh.Style_recodeRegular Cab Pickup < 2e-16 ***
## Veh.Style_recodeSedan 5.42e-05 ***
## Veh.Style_recodeSUV 4.64e-11 ***
## Veh.Style_recodeWagon 5.95e-07 ***
## Engine.Fuel.Typediesel 4.77e-05 ***
## Engine.Fuel.Typeelectric 0.025013 *
## Engine.Fuel.Typeflex-fuel (premium unleaded recommended/E85) 0.001048 **
## Engine.Fuel.Typeflex-fuel (premium unleaded required/E85) 0.000118 ***
## Engine.Fuel.Typeflex-fuel (unleaded/E85) 0.558473
## Engine.Fuel.Typenatural gas 0.026572 *
## Engine.Fuel.Typepremium unleaded (recommended) 0.044975 *
## Engine.Fuel.Typepremium unleaded (required) 0.002420 **
## Engine.Fuel.Typeregular unleaded 0.536046
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##

```



```
par(mfrow=c(2,2))
plot(Model1)
```



```
plot(cooks.distance(Model1))
```

#Cook's D shows outliers but they don't have strong Leverage. 90 % of the residuals are within -2 to +2 and have random cloud and relative constant variance.

```
car::vif(Model1) # Variance Inflation Factor (anything above 10 is a problem)
```

##		GVIF	Df	GVIF^(1/(2*Df))
##	Year	1.275534	1	1.129395
##	Engine.HP	2.678279	1	1.636545
##	Transmission.Type	12.525099	4	1.371585
##	Driven_Wheels	4.038226	3	1.261920
##	Number.of.Doors	58.612240	2	2.766923
##	Vehicle.Size	2.831413	2	1.297182
##	Veh.Style_recode	265.665935	14	1.220628
##	Engine.Fuel.Type	21.813260	9	1.186789

*#there is high multicollinearity with Vehicle style and number of doors.
even though removing it does not have any significant effect on the model,
they should be removed in other to see the actual predictors contributing to
the effect sizes.*

#below the revised model following removal of multicollinear predictors


```
Model1.new<- lm(log(MSRP)~Year+(Engine.HP)^2+Transmission.Type+Driven_Wheels+  
Vehicle.Size+Engine.Fuel.Type, data=car_nonEx.new, )  
summary(Model1.new)
```

```
##  
## Call:  
## lm(formula = log(MSRP) ~ Year + (Engine.HP)^2 + Transmission.Type +  
##     Driven_Wheels + Vehicle.Size + Engine.Fuel.Type, data = car_nonEx.new)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -1.84652 -0.12007 -0.00888  0.11533  1.27863   
##  
## Coefficients:  
##                                     Estimate  
## (Intercept)                      -1.884e+01  
## Year                               1.413e-02  
## Engine.HP                         2.668e-03  
## Transmission.TypeAUTOMATIC        -4.690e-03  
## Transmission.TypeDIRECT_DRIVE      5.050e-02  
## Transmission.TypeMANUAL            -1.596e-01  
## Transmission.TypeUNKNOWN           -1.329e+00  
## Driven_Wheelsfour wheel drive      9.843e-04  
## Driven_Wheelsfront wheel drive     -1.016e-01  
## Driven_Wheelsrear wheel drive      -6.800e-02  
## Vehicle.SizeLarge                  1.043e-01  
## Vehicle.SizeMidsize                5.084e-02  
## Engine.Fuel.Typediesel              5.274e-01  
## Engine.Fuel.Typeelectric            4.035e-01  
## Engine.Fuel.Typeflex-fuel (premium unleaded recommended/E85) 4.345e-01  
## Engine.Fuel.Typeflex-fuel (premium unleaded required/E85)   5.678e-01  
## Engine.Fuel.Typeflex-fuel (unleaded/E85)                     6.163e-02  
## Engine.Fuel.Typenatural gas      4.262e-01  
## Engine.Fuel.Typepremium unleaded (recommended)              2.769e-01  
## Engine.Fuel.Typepremium unleaded (required)                 3.899e-01  
## Engine.Fuel.Typeeregular unleaded 9.928e-02  
##                                     Std. Error t  
value  
## (Intercept)                      1.000e+00 -1  
8.830
```

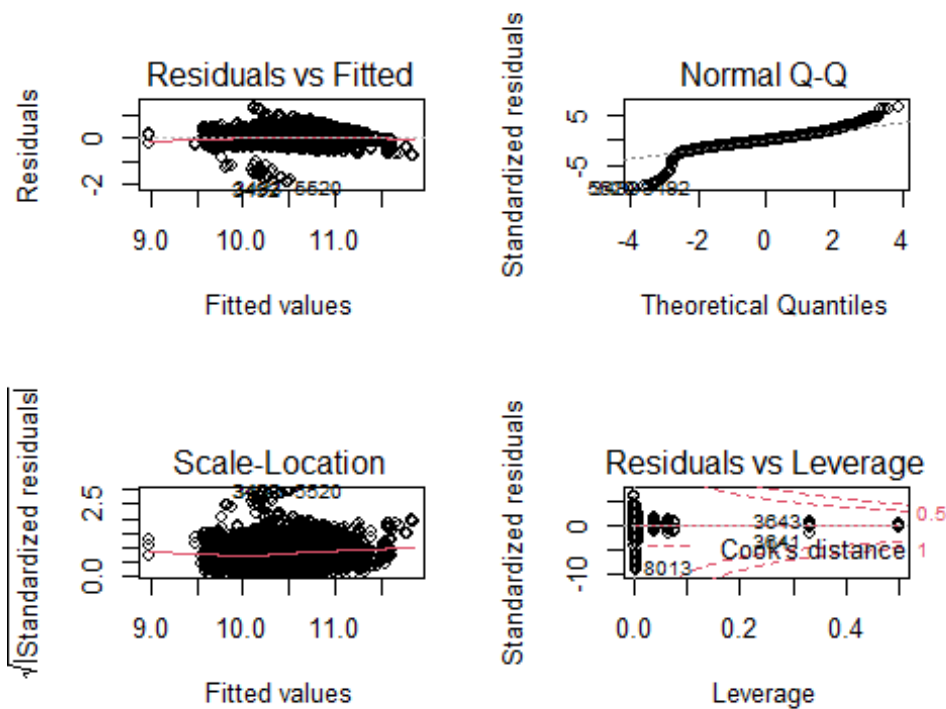
## Year	4.959e-04	2
8.487		
## Engine.HP	3.746e-05	7
1.227		
## Transmission.TypeAUTOMATIC	1.093e-02	-
0.429		
## Transmission.TypeDIRECT_DRIVE	1.459e-01	
0.346		
## Transmission.TypeMANUAL	1.140e-02	-1
4.006		
## Transmission.TypeUNKNOWN	1.196e-01	-1
1.109		
## Driven_Wheelsfour wheel drive	8.492e-03	
0.116		
## Driven_Wheelsfront wheel drive	6.176e-03	-1
6.448		
## Driven_Wheelsrear wheel drive	6.784e-03	-1
0.023		
## Vehicle.SizeLarge	7.493e-03	1
3.927		
## Vehicle.SizeMidsize	5.496e-03	
9.249		
## Engine.Fuel.Typediesel	1.202e-01	
4.388		
## Engine.Fuel.Typeelectric	1.965e-01	
2.054		
## Engine.Fuel.Typeflex-fuel (premium unleaded recommended/E85)	1.256e-01	
3.459		
## Engine.Fuel.Typeflex-fuel (premium unleaded required/E85)	1.298e-01	
4.375		
## Engine.Fuel.Typeflex-fuel (unleaded/E85)	1.192e-01	
0.517		
## Engine.Fuel.Typenatural gas	1.879e-01	
2.268		
## Engine.Fuel.Typepremium unleaded (recommended)	1.191e-01	
2.325		
## Engine.Fuel.Typepremium unleaded (required)	1.191e-01	
3.274		
## Engine.Fuel.Typeregular unleaded	1.189e-01	
0.835		
##	Pr(> t)	
## (Intercept)	< 2e-16	***
## Year	< 2e-16	***
## Engine.HP	< 2e-16	***
## Transmission.TypeAUTOMATIC	0.667903	
## Transmission.TypeDIRECT_DRIVE	0.729253	
## Transmission.TypeMANUAL	< 2e-16	***
## Transmission.TypeUNKNOWN	< 2e-16	***
## Driven_Wheelsfour wheel drive	0.907731	
## Driven_Wheelsfront wheel drive	< 2e-16	***

```
## Driven_Wheelsrear wheel drive < 2e-16 ***
## Vehicle.SizeLarge < 2e-16 ***
## Vehicle.SizeMidsize < 2e-16 ***
## Engine.Fuel.Typediesel 1.16e-05 ***
## Engine.Fuel.Typeelectric 0.040019 *
## Engine.Fuel.Typeflex-fuel (premium unleaded recommended/E85) 0.000544 ***
## Engine.Fuel.Typeflex-fuel (premium unleaded required/E85) 1.23e-05 ***
## Engine.Fuel.Typeflex-fuel (unleaded/E85) 0.605051
## Engine.Fuel.Typenatural gas 0.023356 *
## Engine.Fuel.Typepremium unleaded (recommended) 0.020098 *
## Engine.Fuel.Typepremium unleaded (required) 0.001064 **
## Engine.Fuel.Typeeregular unleaded 0.403728
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2057 on 9550 degrees of freedom
## Multiple R-squared:  0.7582, Adjusted R-squared:  0.7577
## F-statistic: 1497 on 20 and 9550 DF, p-value: < 2.2e-16

par(mfrow=c(2,2))
```



```
plot(Model11.new)
```




```
plot(cooks.distance(Model1.new))
```

```
str(car_nonEx.new)
```

```
## 'data.frame':  9571 obs. of  17 variables:
## $ Make          : Factor w/ 39 levels "Acura","Alfa Romeo",...: 5 5 5 5
5 5 5 5 5 5 ...
## $ Model         : Factor w/ 811 levels "1 Series","1 Series M",...: 2 1
1 1 1 1 1 1 1 1 ...
## $ Year          : int   2011 2011 2011 2011 2011 2012 2012 2012 2012 20
13 ...
## $ Engine.Fuel.Type : Factor w/ 10 levels "", "diesel", "electric",...: 9 9 9
9 9 9 9 9 9 9 ...
## $ Engine.HP      : int   335 300 300 230 230 230 300 300 230 230 ...
## $ Engine.Cylinders : Factor w/ 8 levels "0","3","4","5",...: 5 5 5 5 5 5 5
5 5 5 ...
## $ Transmission.Type: Factor w/ 5 levels "AUTOMATED_MANUAL",...: 4 4 4 4 4
4 4 4 4 4 ...
## $ Driven_Wheels   : Factor w/ 4 levels "all wheel drive",...: 4 4 4 4 4 4
4 4 4 4 ...
## $ Number.of.Doors  : Factor w/ 3 levels "2","3","4": 1 1 1 1 1 1 1 1 1 1
...
## $ Vehicle.Size     : Factor w/ 3 levels "Compact","Large",...: 1 1 1 1 1 1
1 1 1 1 ...
## $ Vehicle.Style    : Factor w/ 16 levels "2dr Hatchback",...: 9 7 9 9 7 9
7 9 7 7 ...
```

```
## $ highway.MPG      : int  26 28 28 28 28 28 26 28 28 27 ...
## $ city.mpg         : int  19 19 20 18 18 18 17 20 18 18 ...
## $ Popularity       : Factor w/ 39 levels "21","26","61",...: 38 38 38 38 3
8 38 38 38 38 38 ...
## $ MSRP             : int  46135 40650 36350 29450 34500 31200 44100 39300
36900 37200 ...
## $ Veh.Style_recode1: Factor w/ 16 levels "2dr Hatchback",...: 9 7 9 9 7 9
7 9 7 7 ...
## $ Veh.Style_recode : Factor w/ 15 levels "2dr Hatchback",...: 7 5 7 7 5 7
5 7 5 5 ...

#Evaluation metrics for the model:
#Root Mean Squared Error
Root_MSE = sqrt(mean(Model1$residuals^2)) #0.19 which is very low and very go
od
Root_MSE_B = sqrt(mean(Model1.new$residuals^2)) #0.205
```




##Observations: 1. The vif output reveals high multicollinearity in Veh.Style_recode, Number.of.Doors and Engine.Fuel.Type predictors. The first two were the most severe (266 and 58 respectively). 2. When these 3 predictors were removed from the model the adjusted Rsquared dropped from 0.78 to 0.69 and the RMSE increased from 0.205 to 0.235. 3. Given this poor performance, I returned the Engine.Fuel.Type predictor given its mildly multicollinear value. Then the adjusted r-squared increased to 0.76 and RMSE shows 0.205 which is similar to what was obtained for the original full model. But now we can explain that the predictors in this model are contributing to the effect size seeing in the response.

#Weighted least squares regression Calculate fitted values from a regression of absolute residuals vs fitted values.

```
library (MASS)
```

```
Weighted_fit <- rlm(log(MSRP)~Year+(Engine.HP)^2+Transmission.Type+Driven_Wheels+Vehicle.Size+Engine.Fuel.Type, data=car_nonEx.new)
```

```
plot(fitted(Weighted_fit), residuals(Weighted_fit))
```



##A bit of non-constant variance here so let use weighted to see if there will be improvement.

```
wts <- 1/fitted(lm(abs(residuals(Weighted_fit)) ~ fitted(Weighted_fit)))^2
```

```
Weighted_fit1 <- rlm(log(MSRP)~Year+(Engine.HP)^2+Transmission.Type+Driven_Wheels+Vehicle.Size+Engine.Fuel.Type, data=car_nonEx.new, weights = wts)
```

```
Root_MSE_B1 = sqrt(mean(Weighted_fit$residuals^2))
```

Transform the MSRP (response) to log for better distribution. It showed better correlation with Engine.HP

```
car_nonEx.newB = car_nonEx.new %>%  
  mutate(MSRP.log = log(MSRP))  
car_nonEx.newB = subset(car_nonEx.newB, select=-c(MSRP))
```

```
car_nonEx.newB = subset(car_nonEx.newB, select=-c(Veh.Style_recode1,MSRP))
```

```
car_nonEx.newB =car_nonEx.newB %>%  
  mutate(Year = as.factor(Year))
```

```
#verify the result  
ncol(car_nonEx.newB)
```

```
## [1] 16
```

```
str(car_nonEx.newB)
```

```
## 'data.frame': 9571 obs. of 16 variables:  
## $ Make : Factor w/ 39 levels "Acura","Alfa Romeo",...: 5 5 5 5  
5 5 5 5 5 5 ...  
## $ Model : Factor w/ 811 levels "1 Series","1 Series M",...: 2 1  
1 1 1 1 1 1 1 1 ...  
## $ Year : Factor w/ 27 levels "1991","1992",...: 21 21 21 21 21  
22 22 22 22 23 ...  
## $ Engine.Fuel.Type : Factor w/ 10 levels "", "diesel", "electric",...: 9 9 9  
9 9 9 9 9 9 9 ...  
## $ Engine.HP : int 335 300 300 230 230 230 300 300 230 230 ...  
## $ Engine.Cylinders : Factor w/ 8 levels "0","3","4","5",...: 5 5 5 5 5 5 5 5  
5 5 5 ...  
## $ Transmission.Type: Factor w/ 5 levels "AUTOMATED_MANUAL",...: 4 4 4 4 4  
4 4 4 4 4 ...  
## $ Driven_Wheels : Factor w/ 4 levels "all wheel drive",...: 4 4 4 4 4 4  
4 4 4 4 ...  
## $ Number.of.Doors : Factor w/ 3 levels "2","3","4": 1 1 1 1 1 1 1 1 1 1  
...  
## $ Vehicle.Size : Factor w/ 3 levels "Compact","Large",...: 1 1 1 1 1 1  
1 1 1 1 ...  
## $ Vehicle.Style : Factor w/ 16 levels "2dr Hatchback",...: 9 7 9 9 7 9  
7 9 7 7 ...  
## $ highway.MPG : int 26 28 28 28 28 28 26 28 28 27 ...  
## $ city.mpg : int 19 19 20 18 18 18 17 20 18 18 ...  
## $ Popularity : Factor w/ 39 levels "21","26","61",...: 38 38 38 38 3  
8 38 38 38 38 38 ...  
## $ Veh.Style_recode : Factor w/ 15 levels "2dr Hatchback",...: 7 5 7 7 5 7  
5 7 5 5 ...  
## $ MSRP.log : num 10.7 10.6 10.5 10.3 10.4 ...
```

```
#Split section for the linear model
```

```
attach(car_nonEx.newB)
```



```
## The following objects are masked from car.new:
##
##      city.mpg, Driven_Wheels, Engine.Cylinders, Engine.Fuel.Type,
##      Engine.HP, highway.MPG, Make, Model, Number.of.Doors, Popularity,
##      Transmission.Type, Vehicle.Size, Vehicle.Style, Year

set.seed(123)
splitPerc = .80
splitPerc2 = .50
trainIndices1 = sample(1:dim(car_nonEx.newB)[1], round(splitPerc * dim(car_nonEx.newB)[1]))
train1 = car_nonEx.newB[trainIndices1,]
test_val1 = car_nonEx.newB[-trainIndices1,]

trainIndices1 = sample(1:dim(test_val1)[1], round(splitPerc2 * dim(test_val1)[1]))
test1 = test_val1[trainIndices1,]
validation1 = test_val1[-trainIndices1,]
dim(car_nonEx.newB)

## [1] 9571   16

dim(train1)

## [1] 7657   16

dim(test1)

## [1] 957   16

dim(validation1)

## [1] 957   16
```

#Fit the regression model using the “train” split. This model used the selected predictors used in the full dataset above

```
model_2 = lm(MSRP.log~(Engine.HP)^2+Transmission.Type+Driven_Wheels+Vehicle.Size+Engine.Fuel.Type+Year, data=train1)
```

```
Root_MSE_2.train = sqrt(mean(model_2$residuals^2)) #0.19 which is very Low and very good
```

```
MSE_2.test = mean((test1$MSRP.log - predict.lm(model_2, test1))^2)
```

```
Root_MSE_2.test = sqrt(MSE_2.test) #This test RMSE is the best technique for evaluating a model. But it Looks high on this model.
```

```
#For weighted Linear Regression
```

```
Weighted_fit <- rlm(log(MSRP)~Year+(Engine.HP)^2+Transmission.Type+Driven_Wheels+Vehicle.Size+Engine.Fuel.Type, data=car_nonEx.new)
wts <- 1/fitted(lm(abs(residuals(Weighted_fit)) ~ fitted(Weighted_fit)))^2

Weighted_fit1 <- rlm(log(MSRP)~Year+(Engine.HP)^2+Transmission.Type+Driven_Wheels+Vehicle.Size+Engine.Fuel.Type, data=car_nonEx.new, weights = wts)
```

Given the huge disparity between Train RMSE and Test RMSE, it seems the model overfit the data. But how come?

#Complex model section

Feature engineering:

Model currently has 811 levels and random forest wants a maximum of 32. It turns out most levels have 1 to 10 observations. These are too few so I will collapse these levels and rename them.

```
#Grouped plot
# Most Levels have Less than 50 observations.
model.count3 = car_nonEx.new %>%
  group_by(Model) %>%
  summarise(count=n())

## `summarise()` ungrouping output (override with `.groups` argument)

attach(model.count3)

## The following object is masked from car_nonEx.newB:
##
##      Model


## The following object is masked from car.new:
##
##      Model

#Most of the model recorded has Less than 25 observations
```


```
model.25 = model.count3 %>% filter(model.count3$count<25)  
model.25$Model = as.character(model.25$Model)
```

#View the new distribution

```
boxplot(model.count3$count, data=model.count3)
```



```
boxplot(model.25$count, data=model.25, main = "model<25")
```



#Recode the model variables that have low counts. Threshold of 25 counts (observations) is used here)

```
car_nonEx.New2 <- car_nonEx.new %>%
  mutate(Model_recode = as.character(Model))

car_nonEx.New2 <- car_nonEx.New2 %>%
  mutate(Model_recode = ifelse(Model_recode %in% as.character(model.25$Model)
, "Model<25", Model_recode)
)

car_nonEx.New2 <- car_nonEx.New2 %>%
  mutate(Model_recode = as.factor(Model_recode)
)
unique(car_nonEx.New2$Model_recode)
```

## [1] Model<25	3 Series	300
## [4] 350Z	370Z	3
## [7] 4 Series	4Runner	500
## [10] 9-3	A3	A4
## [13] Acadia	Accord	Aerio
## [16] ATS Coupe	ATS	B9 Tribeca
## [19] Beetle Convertible	Beetle	C-Class
## [22] Camaro	Camry Solara	Canyon
## [25] CC	Challenger	Charger
## [28] Civic	Colorado	Corolla


```

## [31] Corvette          CR-V          Cruze
## [34] CTS              CX-5          Dakota
## [37] Durango          E-Class      Encore
## [40] Escalade ESV     Escalade      Esteem
## [43] Expedition       Explorer Sport Trac F-150
## [46] Foreza          Forte        Frontier
## [49] Golf GTI         Golf          GTI
## [52] Impreza          Jetta GLI     Jetta SportWagen
## [55] Jetta           Journey       Juke
## [58] Kizashi          MDX           Murano
## [61] MX-5 Miata       New Beetle   Outlander Sport
## [64] Passat          Pathfinder    Pilot
## [67] Q50             Ram Pickup 1500 Range Rover Evoque
## [70] Ranger          S-10         S60
## [73] Sequoia         Sienna       Sierra 1500 Classic
## [76] Sierra 1500     Silverado 1500 Classic Silverado 1500
## [79] Sonata          Sonic        Sorento
## [82] SX4            Tacoma       Terrain
## [85] Tiguan          Titan        TrailBlazer
## [88] Transit Connect Transit Wagon  Traverse
## [91] Tribute         Tundra       Veloster
## [94] Venza          WRX          XC60
## [97] XL-7           XL7          XTS
## 99 Levels: 3 3 Series 300 350Z 370Z 4 Series 4Runner 500 9-3 A3 A4 ... XTS

#model.count3$Model_recode.new = model.count3$Model[model.count3$count < 30] =
"Model<25"

model.25more = model.count3 %>% filter(model.count3$count>25)
boxplot(model.25more$count, data=model.25more, main = "model>25")

```



#Model factor has now been reduced from 810 levels to 99 levels. More work needs to be done though.

#Remove columns that had been recoded

```
car_nonEx.New2 = subset(car_nonEx.New2, select=-c(Veh.Style_recode1, Model, Popularity))
```

#re-order column to make response the last column.

```
car_nonEx.New2b = car_nonEx.New2[,c(1,15,14,2,3,4:13)]
```

#verify the result

```
ncol(car_nonEx.New2b)
```

```
## [1] 15
```

```
str(car_nonEx.New2b)
```

```
## 'data.frame': 9571 obs. of 15 variables:
```

```
## $ Make          : Factor w/ 39 levels "Acura","Alfa Romeo",...: 5 5 5 5 5 5 5 5 5 ...
```

```
## $ Model_recode   : Factor w/ 99 levels "3","3 Series",...: 59 59 59 59 59 59 59 59 59 ...
```

```
## $ Veh.Style_recode : Factor w/ 15 levels "2dr Hatchback",...: 7 5 7 7 5 7 5 7 5 5 ...
```

```
## $ Year           : int  2011 2011 2011 2011 2011 2012 2012 2012 2012 2013 ...
```

```
## $ Engine.Fuel.Type : Factor w/ 10 levels "", "diesel", "electric", ...: 9 9 9
9 9 9 9 9 9 9 ...
## $ Engine.HP       : int   335 300 300 230 230 230 300 300 230 230 ...
## $ Engine.Cylinders : Factor w/ 8 levels "0", "3", "4", "5", ...: 5 5 5 5 5 5 5
5 5 5 ...
## $ Transmission.Type: Factor w/ 5 levels "AUTOMATED_MANUAL", ...: 4 4 4 4 4
4 4 4 4 4 ...
## $ Driven_Wheels    : Factor w/ 4 levels "all wheel drive", ...: 4 4 4 4 4 4
4 4 4 4 4 ...
## $ Number.of.Doors  : Factor w/ 3 levels "2", "3", "4": 1 1 1 1 1 1 1 1 1 1
...
## $ Vehicle.Size     : Factor w/ 3 levels "Compact", "Large", ...: 1 1 1 1 1 1
1 1 1 1 ...
## $ Vehicle.Style    : Factor w/ 16 levels "2dr Hatchback", ...: 9 7 9 9 7 9
7 9 7 7 ...
## $ highway.MPG      : int    26 28 28 28 28 28 26 28 28 27 ...
## $ city.mpg         : int    19 19 20 18 18 18 17 20 18 18 ...
## $ MSRP              : int   46135 40650 36350 29450 34500 31200 44100 39300
36900 37200 ...
```

#Grouped plot for Make

Most Levels have less than 400 observations.

#remove the Make with fewer than 60 observations.

```
make.count = car_nonEx.New2b %>%
```

```
  group_by(Make) %>%
```

```
  summarise(count=n())
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

#filter for Make Level that has less than 80obs

```
make.60 = make.count %>% filter(make.count$count<60)
```

make.60 = as.character(make.60\$Make) #the next set of recode steps worked better with as.character then as.factor steps.

```
make.60less= c("Alfa Romeo", "Aston Martin", "FIAT", "Genesis", "HUMMER", "Lotus", "Maserati", "Oldsmobile", "Plymouth")
```

```
car_nonEx.New2b <- car_nonEx.New2b %>%
```

```
  mutate(Make_recode = as.character(Make))
```

```
car_nonEx.New2b <- car_nonEx.New2b %>%
```

```
  mutate(Make_recode = ifelse(Make_recode %in% make.60less, "Make<60", Make_recode)
```


```
  )
```

```
car_nonEx.New2b <- car_nonEx.New2b %>%
```

```
  mutate(Make_recode = as.factor(Make_recode)
```


```
  )
```

```
boxplot(make.count$count, data=make.count, main = "make")
```



```
boxplot(model.25$count, data=model.25, main = "model<25")
```


model<25



#Ensure MSRP is Last column and remove make since it has now ben recoded.

```
car_nonEx.New3b = subset(car_nonEx.New2b, select=-c(Make))
```

#re-order column to make response the Last column.

```
car_nonEx.New3b = car_nonEx.New3b[,c(15,1,2,3,4:14)]
```

#verify the result

```
ncol(car_nonEx.New3b)
```

```
## [1] 15
```

```
str(car_nonEx.New3b)
```

```
## 'data.frame': 9571 obs. of 15 variables:
```

```
## $ Make_recode : Factor w/ 31 levels "Acura","Audi",...: 3 3 3 3 3 3 3 3 3 3 3 ...
```

```
## $ Model_recode : Factor w/ 99 levels "3","3 Series",...: 59 59 59 59 5 9 59 59 59 59 59 ...
```

```
## $ Veh.Style_recode : Factor w/ 15 levels "2dr Hatchback",...: 7 5 7 7 5 7 5 7 5 5 ...
```

```
## $ Year : int 2011 2011 2011 2011 2011 2012 2012 2012 2012 2013 ...
```

```
## $ Engine.Fuel.Type : Factor w/ 10 levels "", "diesel", "electric",...: 9 9 9 9 9 9 9 9 ...
```

```
## $ Engine.HP : int 335 300 300 230 230 230 300 300 230 230 ...
```

```
## $ Engine.Cylinders : Factor w/ 8 levels "0","3","4","5",...: 5 5 5 5 5 5 5 5 5 5 5 ...
## $ Transmission.Type: Factor w/ 5 levels "AUTOMATED_MANUAL",...: 4 4 4 4 4 4 4 4 4 4 4 ...
## $ Driven_Wheels     : Factor w/ 4 levels "all wheel drive",...: 4 4 4 4 4 4 4 4 4 4 4 ...
## $ Number.of.Doors   : Factor w/ 3 levels "2","3","4": 1 1 1 1 1 1 1 1 1 1 1 ...
## $ Vehicle.Size      : Factor w/ 3 levels "Compact","Large",...: 1 1 1 1 1 1 1 1 1 1 1 ...
## $ Vehicle.Style     : Factor w/ 16 levels "2dr Hatchback",...: 9 7 9 9 7 9 7 9 7 7 ...
## $ highway.MPG       : int   26 28 28 28 28 28 26 28 28 27 ...
## $ city.mpg          : int   19 19 20 18 18 18 17 20 18 18 ...
## $ MSRP              : int  46135 40650 36350 29450 34500 31200 44100 39300 36900 37200 ...
```

#Result: MSRP in last column we have only Model currently as 99 levels that need to be reduced further. I will use Lasso to reduce them.

##adjust the MSRP-HP relationship. ##Scale if needed

```
car_nonEx.New3b = car_nonEx.New3b %>%
  mutate(MSRP.log = log(MSRP)) #Transform MSRP to Log MSRP and remove MSRP
car_nonEx.New3b = subset(car_nonEx.New3b, select=-c(MSRP))
```

VARIABLE SELECTION: LINEAR MODELS

#following lasso result of best lambdas on the full data set, remove Model_recode predictor factor

#This was done backwards. I did lasso on full data set and came back here to remove the variable prior to the train test split. I also did lasso on the train and test sets as well

```
car_nonEx.new4 = subset(car_nonEx.New3b, select=-c(Model_recode)) #not factors with >31 levels
```

```
car_nonEx.new4 = car_nonEx.new4 %>%
```

```
  mutate(Year = as.factor(Year)) #Year was being treated as continuous when it was not. It is an ordinal variable.
```

Par down variables using lasso to select variables

#Split section for complex model using the refined data set.

```
attach(car_nonEx.new4)
```

```
## The following objects are masked from car_nonEx.newB:
```

```
##
```

```
##   city.mpg, Driven_Wheels, Engine.Cylinders, Engine.Fuel.Type,  
##   Engine.HP, highway.MPG, MSRP.log, Number.of.Doors,  
##   Transmission.Type, Veh.Style_recode, Vehicle.Size, Vehicle.Style,  
##   Year
```

```
## The following objects are masked from car.new:
```

```
##
```

```
##   city.mpg, Driven_Wheels, Engine.Cylinders, Engine.Fuel.Type,  
##   Engine.HP, highway.MPG, Number.of.Doors, Transmission.Type,  
##   Vehicle.Size, Vehicle.Style, Year
```

```
set.seed(123)
```

```
splitPerc = .80
```

```
splitPerc2 = .50
```

```
trainIndices = sample(1:dim(car_nonEx.new4)[1], round(splitPerc * dim(car_nonEx.new4)[1]))
```

```
train = car_nonEx.new4[trainIndices,]
```

```
test_val = car_nonEx.new4[-trainIndices,]
```

```
trainIndices1 = sample(1:dim(test_val)[1], round(splitPerc2 * dim(test_val)[1]))
```

```
test = test_val[trainIndices1,]
```

```
validation = test_val[-trainIndices1,]
```

```
dim(car_nonEx.new4)
```

```
## [1] 9571 14
```

```
dim(train)
## [1] 7657  14

dim(test)
## [1] 957  14

dim(validation)
## [1] 957  14
```

Search for the best lambdas used to shrink predictors.

```
#10^10 to 10^-2
grid = 10^seq(10, -2, length = 100)
```

Convert to train-test suitable for lasso regression

```
x_train = model.matrix(MSRP.log~., train)[, -1]
x_test = model.matrix(MSRP.log~., test)[, -1]

y_train = train %>%
  dplyr::select(MSRP.log) %>%
  unlist() %>%
  as.numeric()

y_test = test %>%
  dplyr::select(MSRP.log) %>%
  unlist() %>%
  as.numeric()
```

LAGSSO REGRESSION


Next we fit a lasso regression model on the training set, and evaluate its RMSE on the test set.

We expect the coefficient estimates to be much smaller, in terms of l2 norm, when a large value of λ is used, as compared to when a small value of λ is used.

```
library(glmnet)
lasso_mod = glmnet(x_train,
  y_train,
  alpha = 1,
  lambda = grid) # Fit lasso model on training data

par(mfrow=c(1,1))
plot(lasso_mod) # Draw plot of coefficients

## Warning in regularize.values(x, y, ties, missing(ties), na.rm = na.rm):
## collapsing to unique 'x' values
```



#Snooping on the results of the Lambdas

`lasso_mod$lambda[100]` *#Display 100th Lambda value*

[1] 0.01

`coef(lasso_mod)[,100]` *# Display coefficients associated with 100th Lambda value*

```
## (Intercept)
## 9.715545224
## Make_recodeAudi
## 0.170245545
## Make_recodeBMW
## 0.123164081
## Make_recodeBuick
## 0.000000000
## Make_recodeCadillac
## 0.200234095
## Make_recodeChevrolet
## -0.039874229
## Make_recodeChrysler
## 0.000000000
## Make_recodeDodge
## -0.114220574
## Make_recodeFord
## 0.000000000
```

```

## Make_recodeGMC
## 0.000000000
## Make_recodeHonda
## 0.000000000
## Make_recodeHyundai
## -0.017516740
## Make_recodeInfiniti
## 0.000000000
## Make_recodeKia
## -0.051059103
## Make_recodeLand Rover
## 0.159722893
## Make_recodeLexus
## 0.122905953
## Make_recodeLincoln
## 0.100950287
## Make_recodeMake<60
## 0.108554399
## Make_recodeMazda
## 0.000000000
## Make_recodeMercedes-Benz
## 0.091433017
## Make_recodeMitsubishi
## -0.020484169
## Make_recodeNissan
## -0.024160764
## Make_recodePontiac
## -0.031300109
## Make_recodePorsche
## 0.228824901
## Make_recodeSaab
## 0.054453823
## Make_recodeScion
## -0.007559823
## Make_recodeSubaru
## -0.002854737
## Make_recodeSuzuki
## -0.135110956
## Make_recodeToyota
## 0.000000000
## Make_recodeVolkswagen
## 0.000000000
## Make_recodeVolvo
## 0.126104548
## Veh.Style_recode4dr Hatchback
## -0.003436279
## Veh.Style_recodeCargo Minivan
## 0.000000000
## Veh.Style_recodeCargo Van
## 0.000000000

```

```

##          Veh.Style_recodeConvertible
##          0.124408722
##      Veh.Style_recodeConvertible SUV
##          0.000000000
##          Veh.Style_recodeCoupe
##          0.000000000
##      Veh.Style_recodeCrew Cab Pickup
##          0.000000000
##      Veh.Style_recodeExtended Cab Pickup
##          -0.052415751
##      Veh.Style_recodePassenger Minivan
##          0.001307303
##      Veh.Style_recodePassenger Van
##          0.000000000
##      Veh.Style_recodeRegular Cab Pickup
##          -0.129984649
##          Veh.Style_recodeSedan
##          0.000000000
##          Veh.Style_recodeSUV
##          0.000000000
##          Veh.Style_recodeWagon
##          0.000000000
##          Year1992
##          -0.893822356
##          Year1993
##          -1.165239282
##          Year1994
##          -0.252537998
##          Year1995
##          -1.328415846
##          Year1996
##          -0.801557940
##          Year1997
##          -0.262951977
##          Year1998
##          0.000000000
##          Year1999
##          -0.979689573
##          Year2000
##          -1.387973578
##          Year2001
##          -0.013937202
##          Year2002
##          -0.011181204
##          Year2003
##          0.000000000
##          Year2004
##          0.000000000
##          Year2005
##          0.000000000

```

```

## Year2006
## 0.000000000
## Year2007
## -0.058929983
## Year2008
## -0.005354655
## Year2009
## 0.000000000
## Year2010
## 0.000000000
## Year2011
## 0.000000000
## Year2012
## 0.000000000
## Year2013
## 0.000000000
## Year2014
## 0.000000000
## Year2015
## 0.000000000
## Year2016
## 0.009826981
## Year2017
## 0.003028990
## Engine.Fuel.Typediesel
## 0.177538575
## Engine.Fuel.Typeelectric
## 0.000000000
## Engine.Fuel.Typeflex-fuel (premium unleaded recommended/E85)
## 0.000000000
## Engine.Fuel.Typeflex-fuel (premium unleaded required/E85)
## 0.000000000
## Engine.Fuel.Typeflex-fuel (unleaded/E85)
## -0.051763670
## Engine.Fuel.Typenatural gas
## 0.000000000
## Engine.Fuel.Typepremium unleaded (recommended)
## 0.000000000
## Engine.Fuel.Typepremium unleaded (required)
## 0.038595806
## Engine.Fuel.Typeeregular unleaded
## -0.106918090
## Engine.HP
## 0.002861934
## Engine.Cylinders3
## 0.000000000
## Engine.Cylinders4
## 0.000000000
## Engine.Cylinders5
## 0.000000000

```



```
## Engine.Cylinders6
## 0.009133605
## Engine.Cylinders8
## 0.000000000
## Engine.Cylinders10
## 0.000000000
## Engine.Cylinders12
## 0.000000000
## Transmission.TypeAUTOMATIC
## 0.000000000
## Transmission.TypeDIRECT_DRIVE
## 0.000000000
## Transmission.TypeMANUAL
## -0.122870454
## Transmission.TypeUNKNOWN
## 0.000000000
## Driven_Wheelsfour wheel drive
## 0.000000000
## Driven_Wheelsfront wheel drive
## -0.052816796
## Driven_Wheelsrear wheel drive
## -0.011937064
## Number.of.Doors3
## 0.000000000
## Number.of.Doors4
## 0.000000000
## Vehicle.SizeLarge
## 0.047321511
## Vehicle.SizeMidsize
## 0.000000000
## Vehicle.Style2dr SUV
## 0.000000000
## Vehicle.Style4dr Hatchback
## -0.001632665
## Vehicle.Style4dr SUV
## 0.071305672
## Vehicle.StyleCargo Minivan
## 0.000000000
## Vehicle.StyleCargo Van
## 0.000000000
## Vehicle.StyleConvertible
## 0.005300364
## Vehicle.StyleConvertible SUV
## 0.000000000
## Vehicle.StyleCoupe
## 0.000000000
## Vehicle.StyleCrew Cab Pickup
## 0.000000000
## Vehicle.StyleExtended Cab Pickup
## -0.004516003
```

```
##           Vehicle.StylePassenger Minivan
##                               0.000243327
##           Vehicle.StylePassenger Van
##                               0.000000000
##           Vehicle.StyleRegular Cab Pickup
##                               -0.004961261
##           Vehicle.StyleSedan
##                               0.000000000
##           Vehicle.StyleWagon
##                               0.000000000
##           highway.MPG
##                               0.000000000
##           city.mpg
##                               0.001237789

sqrt(sum(coef(lasso_mod)[-1,100]^2)) # Calculate L1 norm. sqrt of L2 norm, right?


## [1] 2.821179
```

Notice that in the coefficient plot that depending on the choice of tuning parameter, some of the coefficients are exactly equal to zero.

Therefore, we use cross-validation to choose the tuning parameter λ . We can do this using the built-in cross-validation function, `cv.glmnet()`. By default, the function performs 10-fold cross-validation, though this can be changed using the argument `fold`s. Note that we set a random seed first so our results will be reproducible, since the choice of the cross-validation folds is random.

```
set.seed(123)
cv.out = cv.glmnet(x_train, y_train, alpha = 1) # Fit Lasso model on training data

plot(cv.out) # Draw plot of training MSE as a function of Lambda. At the elbow, my best Lambda to minimize MSE is at Log(-5ish)
```



```
bestlam = cv.out$lambda.min # Select lamda that minimizes training MSE
lasso_pred = predict(lasso_mod, s = bestlam, newx = x_test) # Use best Lambda
to predict test data
RMSE.lasso = sqrt(mean((lasso_pred - y_test)^2)) # Calculate test RMSE = #0.1
75

R_Squared = 1 - cv.out$cvm/var(y_train)
max(R_Squared) #= 0.864

## [1] 0.8642786

max(cv.out$glmnet.fit$dev.ratio) #= 0.88. good enough. The model can explain
88% of variation in my dataset. I used corss validation so this should be goo
d with test set.

## [1] 0.8801439

x = model.matrix(MSRP.log~., car_nonEx.new4)[,-15] # trim off the first colum
n
# Leaving only the predictors

y = car_nonEx.new4 %>%
  dplyr::select(MSRP.log) %>%
  unlist() %>%
  as.numeric() #Vector for the target variable.
```

Here we see the number of coefficient estimates are exactly zero:

```
out = glmnet(x, y, alpha = 1, lambda = grid) # Fit Lasso model on full dataset
lasso_coef = predict(out, type = "coefficients", s = bestlam)[1:100,] # Display coefficients using Lambda chosen by CV
```


Selecting only the predictors with non-zero coefficients, we see that the lasso model with λ chosen by cross-validation contains only 54 variables:

```
length(lasso_coef)
## [1] 100

length(lasso_coef[lasso_coef != 0]) # Display only non-zero coefficients
## [1] 54

#When 100 best was given, i got lasso_coef = 54

plot(lasso_coef[lasso_coef != 0]) # Display only non-zero coefficients
```



```
lasso_coef[lasso_coef != 0]

## (Intercept)
## 9.733253872
## Make_recodeAudi
## 0.165285300
## Make_recodeBMW
```

```

##          0.123219297
##      Make_recodeCadillac
##          0.200807909
##      Make_recodeChevrolet
##          -0.042814961
##          Make_recodeDodge
##          -0.114376507
##      Make_recodeHyundai
##          -0.018196217
##          Make_recodeKia
##          -0.047264320
##          Make_recodeLexus
##          0.123369383
##      Make_recodeLincoln
##          0.096809327
##      Make_recodeMake<60
##          0.100785973
##      Make_recodeMercedes-Benz
##          0.082975855
##      Make_recodeMitsubishi
##          -0.025314496
##          Make_recodeNissan
##          -0.024464029
##      Make_recodePontiac
##          -0.035305006
##      Make_recodePorsche
##          0.246710099
##          Make_recodeSaab
##          0.042051657
##          Make_recodeScion
##          -0.013308178
##      Make_recodeSubaru
##          -0.006563362
##      Make_recodeSuzuki
##          -0.134385537
##      Make_recodeVolvo
##          0.124586489
##      Veh.Style_recode4dr Hatchback
##          -0.007187415
##      Veh.Style_recodeConvertible
##          0.113148885
##      Veh.Style_recodeExtended Cab Pickup
##          -0.052587497
##      Veh.Style_recodePassenger Minivan
##          0.003531794
##      Veh.Style_recodeRegular Cab Pickup
##          -0.134250444
##      Veh.Style_recodeSUV
##          0.033009338
##      Year1992

```

```

## -0.826300319
## Year1993
## -1.366907628
## Year1994
## -0.130494509
## Year1995
## -0.926936012
## Year1996
## -0.748808315
## Year1997
## -0.371306719
## Year1998
## -0.129488021
## Year1999
## -1.184716563
## Year2000
## -1.311928880
## Year2001
## -0.023422674
## Year2002
## -0.016141716
## Year2007
## -0.058009798
## Year2008
## -0.005150472
## Year2016
## 0.008379971
## Year2017
## 0.002887764
## Engine.Fuel.Typediesel
## 0.173350815
## Engine.Fuel.Typeflex-fuel (premium unleaded required/E85)
## 0.091588787
## Engine.Fuel.Typeflex-fuel (unleaded/E85)
## -0.061379238
## Engine.Fuel.Typepremium unleaded (required)
## 0.037961584
## Engine.Fuel.Typepremium unleaded
## -0.113344214
## Engine.HP
## 0.002845829
## Engine.Cylinders6
## 0.006996307
## Transmission.TypeMANUAL
## -0.124110760
## Driven_Wheelsfour wheel drive
## 0.006820672
## Driven_Wheelsfront wheel drive
## -0.057641514
## Driven_Wheelsrear wheel drive

```

```

## -0.016253497
## Vehicle.SizeLarge
## 0.043698520

good.lasso = cbind(lasso_coef[lasso_coef != 0])# Display only non-zero coefficients
good.lasso

## [ ,1]
## (Intercept) 9.733253872
## Make_recodeAudi 0.165285300
## Make_recodeBMW 0.123219297
## Make_recodeCadillac 0.200807909
## Make_recodeChevrolet -0.042814961
## Make_recodeDodge -0.114376507
## Make_recodeHyundai -0.018196217
## Make_recodeKia -0.047264320
## Make_recodeLexus 0.123369383
## Make_recodeLincoln 0.096809327
## Make_recodeMake<60 0.100785973
## Make_recodeMercedes-Benz 0.082975855
## Make_recodeMitsubishi -0.025314496
## Make_recodeNissan -0.024464029
## Make_recodePontiac -0.035305006
## Make_recodePorsche 0.246710099
## Make_recodeSaab 0.042051657
## Make_recodeScion -0.013308178
## Make_recodeSubaru -0.006563362
## Make_recodeSuzuki -0.134385537
## Make_recodeVolvo 0.124586489
## Veh.Style_recode4dr Hatchback -0.007187415
## Veh.Style_recodeConvertible 0.113148885
## Veh.Style_recodeExtended Cab Pickup -0.052587497
## Veh.Style_recodePassenger Minivan 0.003531794
## Veh.Style_recodeRegular Cab Pickup -0.134250444
## Veh.Style_recodeSUV 0.033009338
## Year1992 -0.826300319
## Year1993 -1.366907628
## Year1994 -0.130494509
## Year1995 -0.926936012
## Year1996 -0.748808315
## Year1997 -0.371306719
## Year1998 -0.129488021
## Year1999 -1.184716563
## Year2000 -1.311928880
## Year2001 -0.023422674
## Year2002 -0.016141716
## Year2007 -0.058009798
## Year2008 -0.005150472
## Year2016 0.008379971

```

```
## Year2017 0.002887764
## Engine.Fuel.Typediesel 0.173350815
## Engine.Fuel.Typeflex-fuel (premium unleaded required/E85) 0.091588787
## Engine.Fuel.Typeflex-fuel (unleaded/E85) -0.061379238
## Engine.Fuel.Typepremium unleaded (required) 0.037961584
## Engine.Fuel.Typeregular unleaded -0.113344214
## Engine.HP 0.002845829
## Engine.Cylinders6 0.006996307
## Transmission.TypeMANUAL -0.124110760
## Driven_Wheelsfour wheel drive 0.006820672
## Driven_Wheelsfront wheel drive -0.057641514
## Driven_Wheelsrear wheel drive -0.016253497
## Vehicle.SizeLarge 0.043698520
```

#Write the result to a csv file for manipulation and view the result.

```
("C:/Users/olani/OneDrive/Documents/Data Science/SMU-Data Science/Applied Sta  
tistics/Project1Details_2021/Project1Details_2021")
```

```
## [1] "C:/Users/olani/OneDrive/Documents/Data Science/SMU-Data Science/Applied Statistics/Project1Details_2021/Project1Details_2021"
```

```
write.csv(good.lasso, "C:/Users/olani/OneDrive/Documents/Data Science/SMU-Data Science/Applied Statistics/Project1Details_2021/Project1Details_2021/goodlasso.csv")
```

```
goodlassos = read.csv("goodlasso_names.csv")
attach(goodlassos)
str(goodlassos)
```

```
## 'data.frame': 28 obs. of 1 variable:
## $ Names: chr "Make_recodeAudi" "Make_recodeBMW" "Make_recodeCadillac" "M
ake_recodeChevrolet" ...
```

Lasso worked great. it lowered RMSE for regression a lot.

#Following feature engineering. We have a good dataset for tree.

DECISION TREES MODEL

```
library(tree)
```

```
## Warning: package 'tree' was built under R version 4.0.3
```

```
## Registered S3 method overwritten by 'tree':
##   method      from
##   print.tree cli
```



```

set.seed(123)
#Decision tree. We allowed the tree to grow fully without any constraint and
we will prune afterwards.


car.new_tree2 <- tree(MSRP.log~., data=train)

#Let us see the model
summary(car.new_tree2)

##
## Regression tree:
## tree(formula = MSRP.log ~ ., data = train)
## Variables actually used in tree construction:
## [1] "Engine.HP" "Make_recode" "Year"
## Number of terminal nodes: 11
## Residual mean deviance: 0.03795 = 290.2 / 7646
## Distribution of residuals:
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## -2.092000 -0.124500 -0.005558  0.000000  0.122400  1.090000

plot(car.new_tree2)
text(car.new_tree2, pretty=0)

```




#The unpruned tree found "Engine.HP" "Make_recode" and "Year" as the most 'important' variables.

#Let us check if we need to prune the tree?

#We performed cross validation to estimate the cross-validated Mean Squared Error of the trees.


```
cv_car.new_tree2 = cv.tree(car.new_tree2)
plot(cv_car.new_tree2$size, cv_car.new_tree2$dev, type = 'b') #significant best is 10-node tree
```



#Yes we need to prune to the best 10 nodes.

The plot above indicates that 10-node tree is good enough.

```
prune.car.new_tree2 = prune.tree(car.new_tree2, best=10)
plot(prune.car.new_tree2)
text(prune.car.new_tree2, pretty=0)
```




```

#Let's use the pruned tree to make prediction
#prediction
car_pred_new2 <- predict(prune.car.new_tree2, newdata = test)

ggplot() +
  geom_point(aes(x = test$MSRP, y = car_pred_new2)) +
  geom_abline()


```



#The plot above shows the predicted value is correlated well with actual MSRP which is a good thing. But the association is much better for higher MSRP.

`Root_MSE_3 = sqrt(mean((car_pred_new2 - test$MSRP)^2))` *#test RMSE at 9317 which is lower than the one obtained for linear model test RMSE.*

`plot(prune.car.new_tree2)`
`text(prune.car.new_tree2, pretty=0)`



From the tree diagram above. It appears that the Engine.HP value is the most important variable that explains the MSRP of a vehicle. Followed by Make and then Year.

RANDOM FOREST MODEL

#1000 or more decision tree working together whereby majority decision is used to determine the value of the response.

```

library(randomForest)

## Warning: package 'randomForest' was built under R version 4.0.3

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##

## Attaching package: 'randomForest'

```

```
## The following object is masked from 'package:dplyr':
##
##      combine

## The following object is masked from 'package:ggplot2':
##
##      margin

set.seed(123)
car_new.rT <- randomForest(MSRP.log~., data=train, importance =TRUE, mtry=9,
ntree=1000)
```

#The mtry = 9 means we should use 9 predictors for each split of the tree.

Importance=TRUE allows us to view the importance of each variable in the model.

As we can see below Make, Year and Engine.HP account for most of the predicted MSRP of a vehicle.


```
> importance(car_new.rT)
```

	%IncMSE	IncNodePurity
Make_recode	209.56022	248.325918
Veh.Style_recode	63.28708	25.284694
Year	107.76909	109.678191
Engine.Fuel.Type	26.77641	35.136936
Engine.HP	234.63568	739.015771
Engine.Cylinders	14.94020	40.152504
Transmission.Type	81.12897	11.696606
Driven_Wheels	23.58750	16.511576
Number.of.Doors	27.89922	1.440921
Vehicle.Size	62.50075	16.734764
Vehicle.Style	59.20571	29.147054
highway.MPG	60.19749	14.425124
city.mpg	43.71163	15.988468

```
> |
```

```
car_pred <- predict(car_new.rT, newdata = test)
```

```
ggplot() +
  geom_point(aes(x = test$MSRP.log, y = car_pred)) +
  geom_abline()
```



#The plot above offers the best correlation between predicted MSRP and actual MSRP. Looks very strong.

Root_MSE_4 = sqrt(mean((car_pred-test\$MSRP.log)^2)) #test RMSE is half of that from decision tree which is very very nice.

```
summary(car_pred)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  9.302  10.120  10.348  10.385  10.608  11.448
```

TABULATING THE ERRORS

```
RMSE = c(Root_MSE, Root_MSE_B, Root_MSE_B1, Root_MSE_2.train, Root_MSE_2.test, RMSE.lasso, Root_MSE_3, Root_MSE_4)
Model_type = c("Best8_Linear Full", "Best6_Linear Full", "Weighted_Linear Full", "Best6_Linear Train", "Best6_Linear Test", "Lasso_test RMSE", "Single Tree_pruned", "Random Forest")

RMSE_car = cbind(Model_type, RMSE)

RMSE_car
```

##	Model_type	RMSE
## [1,]	"Best8_Linear Full"	"0.193495275586965"
## [2,]	"Best6_Linear Full"	"0.205453021109038"
## [3,]	"Weighted_Linear Full"	"0.206176980068281"
## [4,]	"Best6_Linear Train"	"0.190171422915304"
## [5,]	"Best6_Linear Test"	"0.195255029890895"
## [6,]	"Lasso_test RMSE"	"0.175085056051979"
## [7,]	"Single Tree_pruned Test"	"0.204902058529831"
## [8,]	"Random Forest Test"	"0.117520686511865"

CONCLUSION

Based on the results obtained for the Root Mean Squared. Random Forest produced the lowest value and will generalize much better than the other models even though it scores low on explainability.

John's story ends here